

**SPRINTS**

# EDF Scheduler Implementation in FreeRTOS

---

A Real-time Systems Masterclass Graduation Project

Submitted by:  
Mahmoud Hamdy

Cairo, 2022

## - Contents:

- Introduction.
- EDF Implementation (the 95% changes).
- EDF Implementation (the 5% changes).
- Our real-time example systems.
- Schedulability testing analytically.
- Schedulability testing using simso.
- Schedulability testing using keil.
- Conclusion.
- References.

## - Introduction:

When dealing with real-time systems, there are many scheduling policies to choose from, depending on the application need, hence we make a proper Real-time operating system (RTOS) choice to realize that need. In this context, we are working with FreeRTOS the entire class, whose scheduling policy is very simple and flexible, it's a fixed priority preemptive round-robin scheduler, with some extended abilities to:

- Enable/Disable preemption.
- Enable/Disable time-sharing (RR policy).
- Change task priority in run-time (using APIs).
- Read task state in run-time (using APIs).

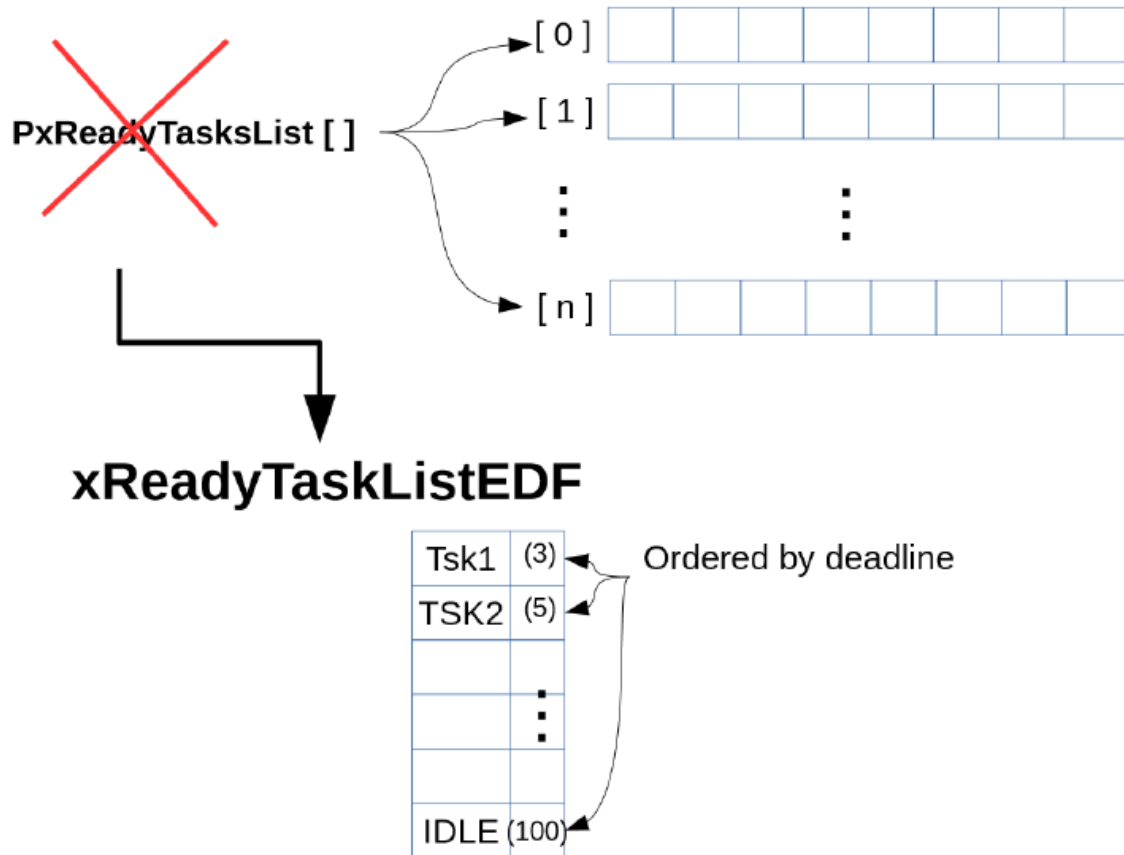
In this project, we try to make good use of the FreeRTOS features and alter its implementation to propose a new type of scheduling policy which is the earliest deadline first (EDF) scheduler, a preemptive scheduler whose task priorities are dynamically assigned - online at every scheduling point - based on deadlines, closer deadline is given higher priority and so on. The new implementation is inspired by a master's degree thesis in computer engineering from University of Padua, more details are mentioned in the "References" section.

I will use this thesis as our theoretical reference for 95% of the implementation and then start introducing new changes one by one (nearly the remaining 5%). Most (if not all) of the changes will be in tasks.c file in the FreeRTOS repository.

We then look at the example systems proposed and perform schedulability testing using three methods: analytically, using simso, and finally on keil simulator with the help of the logic analyzer implemented there. We finally make final comments on the results.

## - EDF Implementation (the 95% changes):

First, a major change is proposed which is how we view the ready task queue, we are going to change it from an array of lists - whose indexes correspond to priority levels - to a list whose node positions correspond to the priority of each task with the head of the list having the highest priority (the closest deadline).



According to FreeRTOS coding standard and style guide, macros are prefixed with the file in which they are defined. The pre-fix is lower case. For example, `configUSE_PREEMPTION` is defined in `FreeRTOSConfig.h`. We use the same approach to create a macro as our configuration variable to enable/disable EDF scheduler, hence `configUSE_EDF_SCHEDULER`. When this macro is set to 1, EDF scheduler is used, else the OS uses the original scheduler.


```
78  #define configUSE_EDF_SCHEDULER 1
```

Then, the next changes are going to happen in `tasks.c`. The new Ready List is declared: `xReadyTasksListEDF` is a simple list structure.

```
361 //EDF code
362 #if (configUSE_EDF_SCHEDULER == 1)
363 #define IDLE_PERIOD (TickType_t)100
364 PRIVILEGED_DATA static List_t xReadyTasksListEDF;
365 #endif
```

Then, `prvInitialiseTaskLists` method, that initializes all the task lists at the creation of the first task, is modified adding the initialization of `xReadyTasksListEDF`

```
3827 static void prvInitialiseTaskLists( void )
3828 {
3829     UBaseType_t uxPriority;
3830
3831     for( uxPriority = ( UBaseType_t ) 0U; uxPriority < ( UBaseType_t ) configMAX_PRIORITIES; uxPriority++ )
3832     {
3833         vListInitialise( &(amp; pxReadyTasksLists[ uxPriority ]) );
3834     }
3835
3836 //EDF code
3837 #if (configUSE_EDF_SCHEDULER == 1)
3838     vListInitialise(&xReadyTasksListEDF);
3839 #endif
3840
3841     vListInitialise( &xDelayedTaskList1 );
3842     vListInitialise( &xDelayedTaskList2 );
3843     vListInitialise( &xPendingReadyList );
3844 }
```



`prvAddTaskToReadyList` method that adds a task to the Ready List is then modified as follows:

```
220 //EDF code: prvAddTaskToReadyList
221 #if (configUSE_EDF_SCHEDULER == 0)
222 #define prvAddTaskToReadyList( pxTCB )
223     traceMOVED_TASK_TO_READY_STATE( pxTCB );
224     taskRECORD_READY_PRIORITY( ( pxTCB )->uxPriority );
225     vListInsertEnd( &(amp; pxReadyTasksLists[ ( pxTCB )->uxPriority ] ), &( ( pxTCB )->xStateListItem ) );
226     tracePOST_MOVED_TASK_TO_READY_STATE( pxTCB );
227 #else
228 #define prvAddTaskToReadyList( pxTCB ) /* xStateListItem must contain the deadline value */
229     traceMOVED_TASK_TO_READY_STATE( pxTCB );
230     vListInsert( &xReadyTasksListEDF, &( ( pxTCB )->xStateListItem ) );
231     tracePOST_MOVED_TASK_TO_READY_STATE( pxTCB );
232 #endif
```

vListInsert() method is called to insert in xReadyTasksListEDF the task TCB pointer. The item will be inserted into the list in a position determined by its item value xStateListItem (ascending item value order). So, it is assumed that xStateListItem contains the next task deadline.

Speaking of deadlines, time for a change in task structure, when a task moves to the Ready List, the knowledge of its next deadline is needed to insert it in the correct position. The deadline is calculated as:

Next deadline = current tick + task period

So, every task needs to store its period value. A new variable is added in the tskTaskControlBlock structure (TCB):

```
263 typedef struct tskTaskControlBlock /* The old naming convention is
264 {
265     volatile StackType_t * pxTopOfStack; /*< Points to the location of th
266
267     #if ( portUSING_MPU_WRAPPERS == 1 )
268         xMPU_SETTINGS xMPUSettings; /*< The MPU settings are defined as p
269     #endif
270
271     // -EDF- code: Period value to help in task deadline calculation
272     #if (configUSE_EDF_SCHEDULER == 1)
273         TickType_t xTaskPeriod; /*< Stores the period in tick of the task */
274     #endif
275
276     ListItem_t xStateListItem; /*< The list that the sta
277     ListItem_t xEventListItem; /*< Used to reference a t
278     UBaseType_t uxPriority; /*< The priority of the t
279     StackType_t * pxStack; /*< Points to the start o
280     char pcTaskName[ configMAX_TASK_NAME_LEN ]; /*< Descriptive name give
```

Accordingly, a new initialization task method is created. `xTaskPeriodicCreate()` is a modified version of the standard method `xTaskGenericCreate()`, that receives the task period as additional input parameter and set the `xTaskPeriod` variable in the task TCB structure.

```

838 //EDF-code: xTaskPeriodicCreate()
839 #if (configUSE_EDF_SCHEDULER == 1)
840 BaseType_t xTaskPeriodicCreate( TaskFunction_t pxTaskCode,
841                                const char * const pcName, /*lint !e971 Unqualified char
842                                const configSTACK_DEPTH_TYPE usStackDepth,
843                                void * const pvParameters,
844                                UBaseType_t uxPriority,
845                                TaskHandle_t * const pxCreatedTask, TickType_t period )
846 {
847     TCB_t * pxNewTCB;
848     BaseType_t xReturn;

```

Before adding the new task to the Ready List by calling `prvAddTaskToReadyList()`, the task's `xStateListItem` is initialized to the value of the next task deadline.

```

918 //EDF-code:
919 pxNewTCB->xTaskPeriod = period; /* Initialize the period */
920 /* Insert the period value in the xStateListItem before adding task to RL */
921 listSET_LIST_ITEM_VALUE(&(pxNewTCB->xStateListItem), pxNewTCB->xTaskPeriod + xTaskGetTickCount());
922
923 prvAddNewTaskToReadyList( pxNewTCB );
924 xReturn = pdPASS;
925 }
926 else
927 {
928     xReturn = errCOULD_NOT_ALLOCATE_REQUIRED_MEMORY;
929 }
930
931 return xReturn;
932 }
933 #endif /* xTaskPeriodicCreate() */

```

The IDLE task management is modified as well. The initialization of the IDLE task happens in the `vTaskStartScheduler()` method, that starts the real time kernel tick processing and initialize all the scheduler structures. Since FreeRTOS specifications want a task in execution at every instant, a correct management of the IDLE task is fundamental. With the standard FreeRTOS scheduler, the IDLE task is a simple task initialized at the lowest priority. In this way it would be scheduled only when no other tasks are in the ready state. With the EDF scheduler, the lowest priority behaviour can be simulated by a task having the furthest deadline. `vTaskStartScheduler()` method initializes the IDLE task and inserts it into the Ready List. The method is modified as follows:

```

2153  ▾  #else /* if ( configSUPPORT_STATIC_ALLOCATION == 1 ) */
2154  ▾  {
2155      /* The Idle task is being created using dynamically
2156
2157  ▾  // EDF code: Initialize IDLE task
2158
2159      #if (configUSE_EDF_SCHEDULER == 1)
2160  ▾  xReturn = xTaskPeriodicCreate( prvIdleTask,
2161      configIDLE_TASK_NAME,
2162      configMINIMAL_STACK_SIZE,
2163      ( void * ) NULL,
2164      portPRIVILEGE_BIT, /* In effect ( tskI
2165      &xIdleTaskHandle, /*lint !e961 MISRA ex
2166      IDLE_PERIOD);
2167
2168  ▾  #else
2169      xReturn = xTaskCreate( prvIdleTask,
2170      configIDLE_TASK_NAME,
2171      configMINIMAL_STACK_SIZE,
2172      ( void * ) NULL,
2173      portPRIVILEGE_BIT, /* In ef
2174      &xIdleTaskHandle ); /*lint !
2175  ▾  #endif
2176  ▾  }

```

The IDLE task is initialized with a period of `IDLE_PERIOD = 100`. We assume that no task can have a period greater than `IDLE_PERIOD`: in this way, when the IDLE task is added to the Ready List, it will be at the last position of the list, since its deadline will be greater than any other task (Next deadline = current tick + task period, with current tick = 0 and task period = `IDLE_PERIOD = 100`, which is greater than any other task period). Every time IDLE task executes (i.e., no other tasks are in the Ready List), it calls a method that increments its deadline to guarantee that IDLE task will remain in the last position of the Ready List. We cover this part in the next section (the 5% changes).



Last change needed involves the switch context mechanism. Every time the running task is suspended, or a suspended task with a higher priority than the running task awakes, a switch context occurs. `vTaskSwitchContext()` method is in charge to update the `pxCurrentTCB`, pointer to the new running task:

```
3177 void vTaskSwitchContext( void )
3178 {
```

```
3228  /* Select a new task to run using either the generic C or port
3229   * optimised asm code. */
3230
3231  /*-EDF-code: Context switching to the earliest deadline task
3232   */
3232  #if (configUSE_EDF_SCHEDULER == 0)
3233      taskSELECT_HIGHEST_PRIORITY_TASK(); /*lint !e9079 void * is used as this ma
3234  #else
3235      pxCurrentTCB = (TCB_t *) listGET_OWNER_OF_HEAD_ENTRY(&xReadyTasksListEDF);
3236  #endif
3237
```

```
3257 }
3258 /*-----
3259
```

`taskSELECT_HIGHEST_PRIORITY_TASK()` method is replaced in order to assign to `pxCurrentTCB` the task at the first place of the new Ready List.

Now we have 95% of the pieces to get the new EDF scheduler work. In the next section will be looking at the remaining 5% of the changes which are not shown in the thesis but a must to have a fully implemented EDF scheduler.

## - EDF Implementation: the 5% changes

First, when we create a new task, its TCB goes through some steps before being added to the Ready List, one of which is checking if the scheduler is not currently running, which often happens at the beginning system initialization. We check if this task has a higher priority (earlier deadline) than the current task and we update the current task based on this result but it has to be a comparison of deadlines not priorities, we add this change in prvAddNewTaskToReadyList():

```
1191 static void prvAddNewTaskToReadyList( TCB_t * pxNewTCB )
1192 {
```

```
1219     /* If the scheduler is not already running, make this task the
1220      * current task if it is the highest priority task to be created
1221      * so far. */
1222     if( xSchedulerRunning == pdFALSE )
1223     {
1224         // EDF code: select current task if scheduler is not running
1225         #if (configUSE_EDF_SCHEDULER == 0)
1226             if( pxCurrentTCB->uxPriority <= pxNewTCB->uxPriority )
1227             {
1228                 pxCurrentTCB = pxNewTCB;
1229             }
1230             else
1231             {
1232                 mtCOVERAGE_TEST_MARKER();
1233             }
1234         #else
1235             if( pxCurrentTCB->xStateListItem.xItemValue > pxNewTCB->xStateListItem.xItemValue )
1236             {
1237                 pxCurrentTCB = pxNewTCB;
1238             }
1239             else
1240             {
1241                 mtCOVERAGE_TEST_MARKER();
1242             }
1243         #endif
1244     }
```

```
1284 }
1285 /*-----*/
```

We must also update the deadline for the idle task each time it executes so that we can always keep it at the lowest priority:

```
3597 static portTASK_FUNCTION( prvIdleTask, pvParameters )
3598 {
3599     /* Stop warnings. */
3600     ( void ) pvParameters;
3601
3602     /** THIS IS THE RTOS IDLE TASK - WHICH IS CREATED AUTOMATICALLY WHEN
3603      * SCHEDULER IS STARTED. */
3604
3605     /* In case a task that has a secure context deletes itself, in which
3606      * the idle task is responsible for deleting the task's secure context
3607      * any. */
3608     portALLOCATE_SECURE_CONTEXT( configMINIMAL_SECURE_STACK_SIZE );
3609
3610     for( ; ; )
3611     {
3612         /*-EDF-code: Update Idle Task Deadline
3613         → pxCurrentTCB->xStateListItem.xItemValue += (TickType_t)100;
3614
3615         /* See if any tasks have deleted themselves - if so then the idle
3616          * is responsible for freeing the deleted task's TCB and stack. */
3617         prvCheckTasksWaitingTermination();
```

Finally, one of the most crucial changes is when a task unblocks (moves from delayed list to ready list), we must check if there should be a context switching based on our new scheduler. These changes are added to the xTaskIncrementTick() function:

```
2873 BaseType_t xTaskIncrementTick( void )
2874 {
2875     TCB_t * pxTCB;
2876     TickType_t xItemValue;
2877     BaseType_t xSwitchRequired = pdFALSE;
```

We first update the task's new deadline after unblocking and before adding to the ready list:

```
2959 //EDF-code: update task new deadline before adding to ready list
2960 #if (configUSE_EDF_SCHEDULER == 1)
2961     listSET_LIST_ITEM_VALUE(&(pxTCB->xStateListItem), pxTCB->xTaskPeriod + listGET_LIST_ITEM_VALUE(&(pxTCB->xStateListItem)));
2962     //pxTCB->xStateListItem.xItemValue = xItemValue + pxTCB->xTaskPeriod;
2963 #endif
2964
2965 /* Place the unblocked task into the appropriate ready
2966  * list. */
2967 prvAddTaskToReadyList( pxTCB );
```

Then we check (based on how our new Ready List looks like after adding the new task) if there should be a context switching by comparing with the current task's deadline (determining the new currentTCB based on currentTCB's deadline):

```
2965 v /* Place the unblocked task into the appropriate ready
2966  * list. */
2967     prvAddTaskToReadyList( pxTCB );
2968
2969 v /* A task being unblocked cannot cause an immediate
2970  * context switch if preemption is turned off. */
2971 v #if ( configUSE_PREEMPTION == 1 )
2972 v {
2973 v     /* Preemption is on, but a context switch should
2974  * only be performed if the unblocked task has a
2975  * priority that is equal to or higher than the
2976  * currently executing task. */
2977
```

```
2978 //EDF-code: Preemption decision based on deadline not priority
2979 #if (configUSE_EDF_SCHEDULER == 1)
2980     if( listGET_LIST_ITEM_VALUE(&(pxTCB->xStateListItem)) <= listGET_LIST_ITEM_VALUE(&(pxCurrentTCB->xStateListItem)) )
2981     {
2982         xSwitchRequired = pdTRUE;
2983     }
2984     else
2985     {
2986         mtCOVERAGE_TEST_MARKER();
2987     }
2988 #else
2989     if( pxTCB->uxPriority >= pxCurrentTCB->uxPriority )
2990     {
2991         xSwitchRequired = pdTRUE;
2992     }
2993     else
2994     {
2995         mtCOVERAGE_TEST_MARKER();
2996     }
2997 #endif
2998 }
2999 #endif /* configUSE_PREEMPTION */
```

It's also worth mentioning that a small change is required in list.c:

```
115 void vListInsert( List_t * const pxList,
116 | | | | | ListItem_t * const pxNewListItem )
117 {
```

```
163 // EDF code: Fine adjustments to the EDF preemption according to the thesis requirements
164 #if(configUSE_EDF_SCHEDULER == 1)
165 for( pxIterator = ( ListItem_t * ) &( pxList->xListEnd ); pxIterator->pxNext->xItemValue < xValueOfInsertion; pxIterator = pxIterator->pxNext )
166 {
167     /* There is nothing to do here, just iterating to the wanted
168     | * insertion position. */
169 }
170 #else
171 for( pxIterator = ( ListItem_t * ) &( pxList->xListEnd ); pxIterator->pxNext->xItemValue <= xValueOfInsertion; pxIterator = pxIterator->pxNext )
172 {
173     /* There is nothing to do here, just iterating to the wanted
174     | * insertion position. */
175 }
176 #endif
```

## - Our real-time example systems

System 1:

|        | $T$ | $C$ |
|--------|-----|-----|
| Task A | 5   | 2   |
| Task B | 8   | 2   |

System 2:

|        | $T$ | $C$ |
|--------|-----|-----|
| Task A | 5   | 3   |
| Task B | 8   | 3   |

Tick time for both systems = H.C.F(5, 8) = 1 ms

Hyper-period for both systems = L.C.M(5, 8) = 5 x 8 = 40 ms

## - Schedulability testing analytically:

From thesis:


**Theorem 3.1** *A task set of periodic tasks is schedulable by EDF if and only if:*

$$U = \sum_{i=1}^N \frac{C_i}{T_i} \leq 1$$

**System 1:**

$$U = C_1/T_1 + C_2/T_2 = 2/5 + 2/8 = 16/40 + 10/40 = 26/40 = 65\% = 0.65 < 1$$


Hence, system is healthy and schedulable by EDF according to the theorem.



**System 2:**


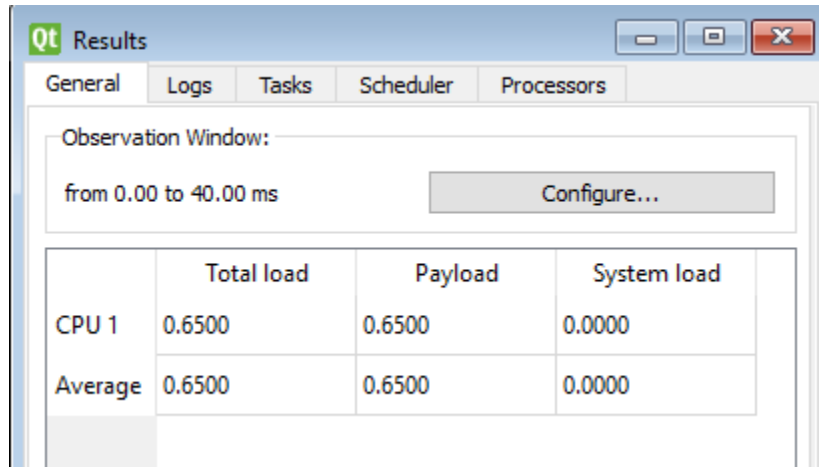
$$U = C_1/T_1 + C_2/T_2 = 3/5 + 3/8 = 24/40 + 15/40 = 39/40 = 97.5\% = 0.975 < 1$$

Hence, system is healthy and schedulable by EDF according to the theorem.




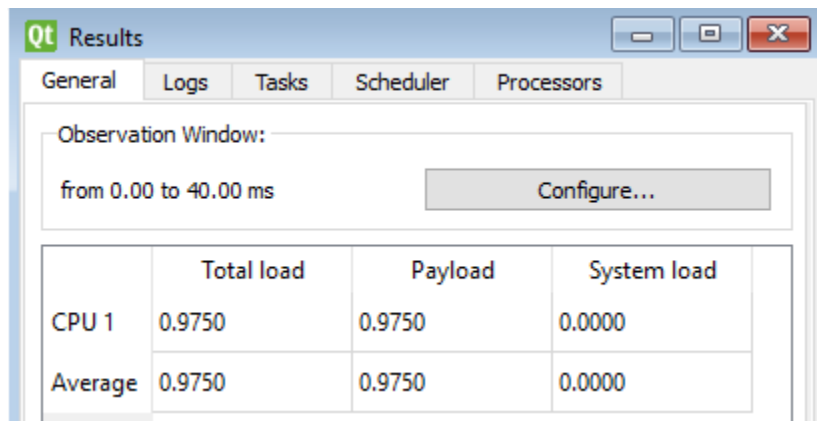
## - Schedulability testing using simso:

### System 1:



Results on simso match the analytical results.

## System 2:



Results on simso (using EDF2) match the analytic results.



## - Schedulability testing using keil:

In this simulation we use GPIO pins with our system hooks using the following configuration:

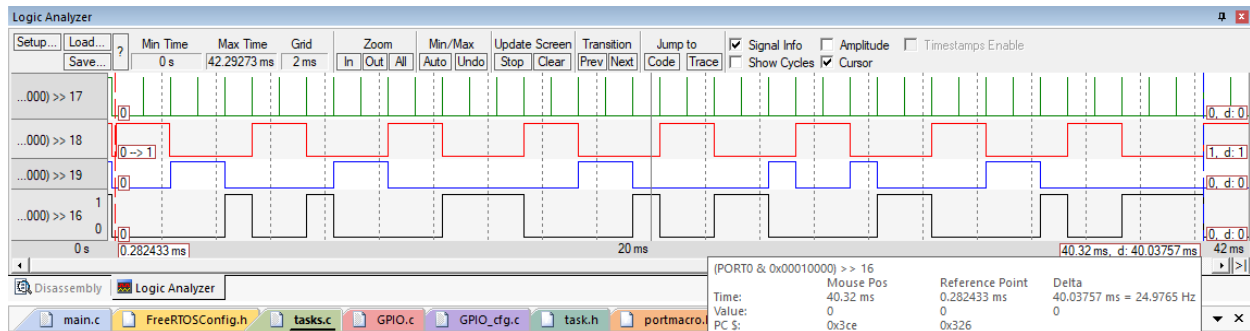
Port 0 pin 17 (PIN1) → Refers to the system tick

Port 0 pin 18 (PIN2) → Refers to Task A

Port 0 pin 19 (PIN3) → Refers to Task B

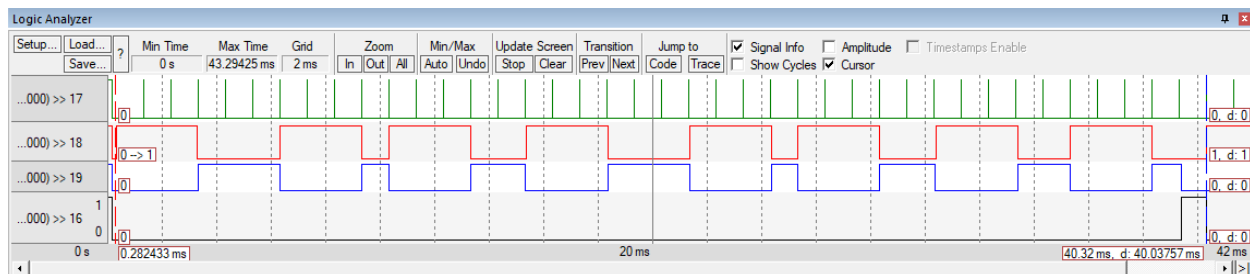
Port 0 pin 16 (PIN0) → Refers to the IDLE Task

### System 1:



Results on keil match simso and the analytical.

### System 2:



Results on keil match simso and the analytical.

## - Conclusion:

According to the results obtained, we can safely say that system 1 and system 2 are both healthy (with one more healthy than the other) and schedulable in a dynamic scheduling policy namely the EDF scheduler and this was proved analytically using the theorem of EDF scheduling, using a simulator like Simso, and finally using real-time code on keil environment with a powerful software simulator and logic analyzer.

## - References:

Carraro, Enrico (2016), Implementation and Test of EDF and LLREF Schedulers in FreeRTOS, Padua @ thesis, accessed 20 July 2021,

[http://tesi.cab.unipd.it/51896/1/Implementation and Test of EDF and LLREF Scgheduler in FreeRTOS.pdf](http://tesi.cab.unipd.it/51896/1/Implementation_and_Test_of_EDF_and_LLREF_Scgheduler_in_FreeRTOS.pdf)

FreeRTOS, accessed 1 January 2022, <https://www.freertos.org/FreeRTOS-Coding-Standard-and-Style-Guide.html>